Capstone Project: Breast Cancer Diagnostic Analysis

Prepared by: Arnav Gehlot

Section 1: Data Overview and Executive Summary

Dataset Selected: Breast Cancer Wisconsin (Diagnostic) Dataset (Source: UCI Machine Learning Repository)

Data Summary:

The dataset contains structured mathematical measurements derived from digitized images of fine needle aspirates (FNA) of breast masses.

- **Size:** Following initial loading, the dataset contained 569 instances (rows) and 33 attributes (columns).
- **Variables:** The features consist of continuous numerical variables representing cellular characteristics computed for each cell nucleus, including metrics such as radius, texture, perimeter, area, and smoothness.
- **Target Variable:** The diagnosis column is the target variable for future classification modeling. It is a categorical variable with two classes: 'M' (Malignant) and 'B' (Benign).

Section 2: Data Exploration Plan

To ensure a rigorous and actionable analysis, the exploration followed this structured plan:

1. **Data Cleaning:** Identify and remove non-predictive identifiers and hidden missing values or formatting artifacts to ensure model integrity.
2. **Target Variable Analysis:** Evaluate the distribution of the 'diagnosis' class to establish baseline accuracy and identify potential class imbalance.
3. **Univariate Feature Distribution:** Utilize boxplots to visually evaluate how individual continuous variables (e.g., cell radius) differ between benign and malignant cases.
4. **Bivariate Correlation Analysis:** Utilize scatter plots to observe the interaction between multiple features simultaneously (e.g., radius vs. texture) to justify the need for multivariate machine learning algorithms.

Section 3: Exploratory Data Analysis (EDA) Results

1. Target Variable Distribution:

Our initial analysis reveals that the dataset is mildly imbalanced. Out of 569 patient records, 357 (62.7%) are diagnosed as Benign, and 212 (37.3%) are diagnosed as Malignant. Any future predictive model must establish a baseline accuracy greater than 62.7% to be considered effective.

2. Univariate Analysis: Cell Radius vs. Diagnosis

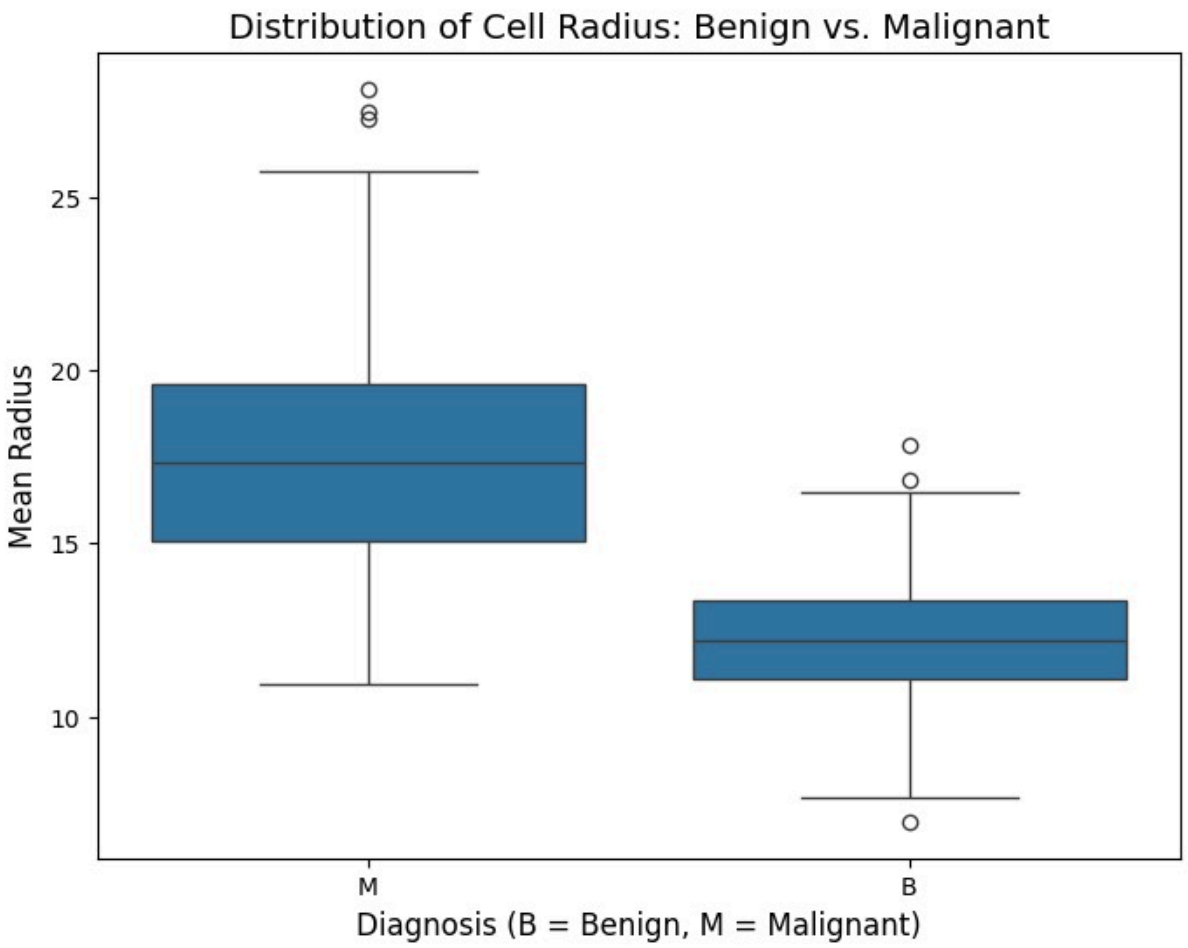

A comparative analysis of the radius_mean feature indicates a strong visual distinction between the two classes. The median radius for malignant cells is notably higher (approx. 17.5) than that of benign cells (approx. 12). Furthermore, the interquartile ranges barely overlap, suggesting high predictive power.

3. Bivariate Analysis: Cell Radius vs. Texture

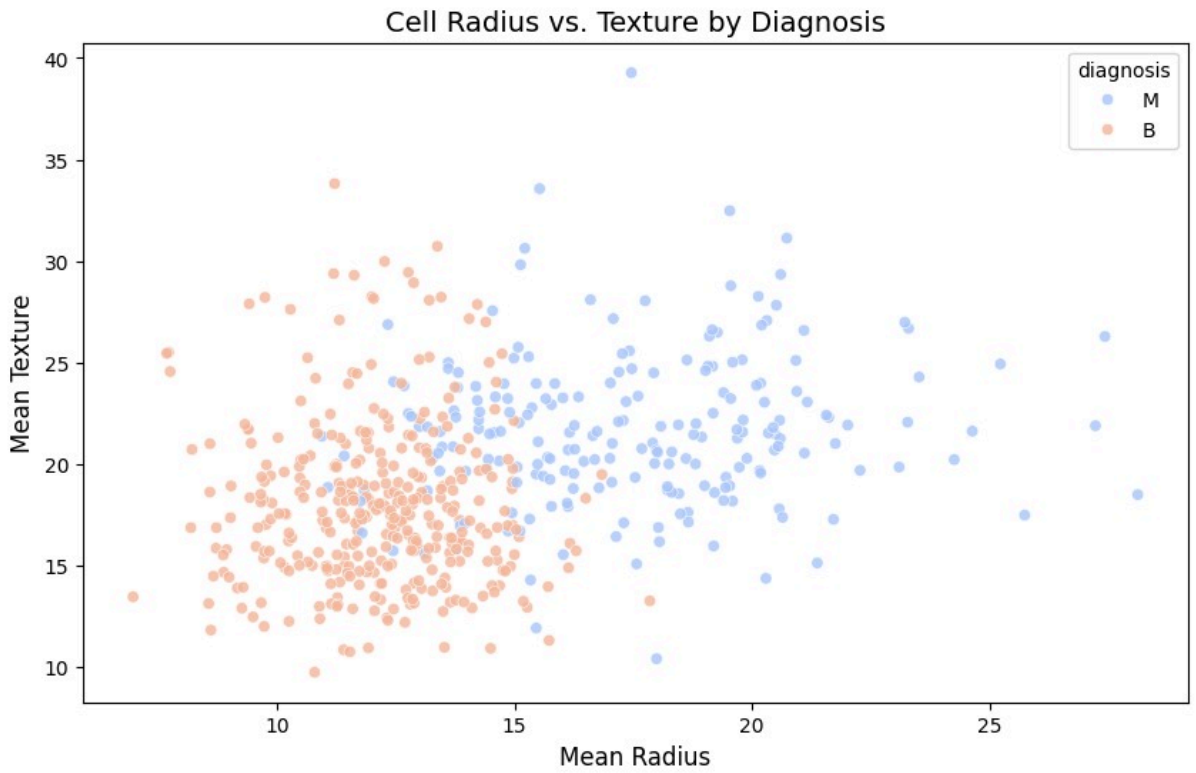

Plotting radius against texture reveals distinct clustering. Benign cases aggregate in the lower-left quadrant (smaller and smoother), while malignant cases extend into the upper-right (larger and rougher). However, a central "overlap zone" exists, proving that a single variable is insufficient for perfect classification.

Section 4: Data Cleaning & Feature Engineering Steps

To prepare the dataset for algorithmic processing, structural formatting issues were addressed:

- **Handling Identifiers:** The id column was dropped entirely. Patient tracking numbers possess no mathematical relationship to biological malignancy and would introduce artificial bias into a machine learning model.
- **Handling Missing Values/Artifacts:** A programmatic check identified and dropped entirely empty columns (e.g., Unnamed: 32), a common artifact in downloaded CSV files.

- **Output:** The dataset was successfully reduced from (569, 33) to a clean, model-ready shape of (569, 31).

Section 5: Key Findings & Insights

The exploratory phase yielded several actionable insights for the predictive modeling phase:

1. **Feature Importance:** Mathematical measurements of the cell nucleus, particularly size (radius_mean) and surface variation (texture_mean), are strong indicators of malignancy. These features should be prioritized.
2. **The Necessity of Machine Learning:** While distinct clusters exist, the visual overlap between benign and malignant profiles in the mid-ranges confirms that simple threshold rules (e.g., "if radius > 15, then malignant") will result in false negatives. A multidimensional model analyzing all 30 variables simultaneously is required for medical-grade accuracy.

Section 6: Hypothesis Formulation

Based on the visual evidence gathered during EDA, the following hypotheses were formulated for statistical validation:

1. **Hypothesis 1 (Radius):**
    - The Null Hypothesis : There is no statistically significant difference in the mean radius between benign and malignant cells.
    - The Alternative Hypothesis : Malignant cells have a significantly larger mean radius than benign cells.
2. **Hypothesis 2 (Texture):**
    - The Null Hypothesis : There is no statistically significant difference in the mean texture between benign and malignant cells.
    - The Alternative Hypothesis : Malignant cells exhibit a significantly higher mean texture measurement than benign cells.
3. **Hypothesis 3 (Correlation):**
    - The Null Hypothesis : There is no linear relationship between a cell's radius and its texture.
    - The Alternative Hypothesis : There is a positive correlation between cell radius and cell texture.

Section 7: Significance Testing

An independent two-sample t-test was conducted to mathematically validate Hypothesis 1 (Cell Radius), using a standard alpha level of 0.05.

- **Test Metric:** radius_mean (Malignant vs. Benign)
- **T-Statistic:** 25.4358
- **P-Value:** 8.46e-96

Insight and Interpretation:

Because the p-value is practically zero (and far below our 0.05 threshold), we firmly reject the null hypothesis. The mathematical probability that the difference in size between these cells is due to random chance is negligible.

Business Impact: We can confidently state to medical stakeholders that radius_mean is a mathematically sound, highly reliable feature for predicting malignancy.

Section 8: Conclusion and Next Steps

This Exploratory Data Analysis successfully evaluated the dataset and confirmed its readiness for predictive modeling. We established with statistical certainty that key numerical features possess strong predictive power in distinguishing between benign and malignant cell nuclei.

Having verified the quality of this dataset, the immediate next steps are:

1. **Feature Scaling:** Apply standardization to the numerical features (since measurements like 'area' and 'smoothness' operate on vastly different mathematical scales) to ensure they are properly weighted by machine learning algorithms.
2. **Model Selection:** Implement supervised machine learning classification models. Given the clear mathematical boundaries observed, algorithms such as Logistic Regression, Support Vector Machines (SVM), and Random Forests will be tested.
3. **Model Evaluation:** Train the models using a standard train/test split and evaluate their performance using strict medical-grade metrics, prioritizing Recall (Sensitivity) to aggressively minimize the risk of false-negative diagnoses.